

# CPPServer



High-performance C++  
microservices for Kubernetes  
and Cloud

## Contact

<https://cppserver.com>

**Email:**  
[cppserver@martincordova.com](mailto:cppserver@martincordova.com)

Martín Córdova y Asociados, C.A.  
Caracas – Venezuela

## Anatomy of a C++ microservice

Our CPPServer platform allows fast development of microservices that return JSON datasets from SQL databases without writing any code, just by declaring tasks using a simple JSON configuration file.



CPPServer is distributed as a docker image to be run as a container on enterprise-grade orchestration platforms, like Kubernetes or Cloud serverless services like Azure Container Apps or AWS Beanstalk.

Documentation available on GitHub:

<https://github.com/cppservergit/cppserver-docs>

Docker Hub repository:

<https://hub.docker.com/r/cppserver/pgsql>

It can also be used as a Linux native *systemd* service. X86-64 and Arch-64 hardware architectures are supported.

## The SQL backend

---

In this particular case we are going to export a multi-array response, given a customer ID the microservice will return a customer dataset and the customer's orders dataset, all in one roundtrip.

<https://cppserver.com/ms/customer/info?customerid=BERGS>

We will use SQLServer as the database for this example, and a stored procedure to return multiple resultsets in one roundtrip to the database server, looking for efficiency, pre-compiled optimized queries and conscious usage of network resources by returning all data in one shot, TransactSQL makes this task very easy to accomplish.

This stored procedure will return two resultsets given a customer ID. Our database backend is ready and follows best practices, we are only returning the columns required to reduce the size of the data packets to be transferred over the network. Tables are properly indexed and statistics are up to date. So far so good.

```
CREATE OR ALTER procedure [dbo].[sp_getCustomerInfo](@customerid varchar(10)) as
begin

    set nocount on

    select
        customerid, contactname, companyname, city, country, phone
    from
        demo..customers
    where customerid = @customerid;

    select orders.orderid,
           orderdate,
           shipcountry,
           shippers.companyname as shipper,
           total
    from
        orders, shippers, vw_order_totals
    where
        customerid = @customerid
    and
        shippers.shipperid = orders.shipvia
    and
        vw_order_totals.orderid = orders.orderid
    order by orderid

end
GO
```

## JSON output

---

According to CPPServer JSON specification, if there is no error processing the request, the output will look like this:

```
{
  "status": "OK",
  "data": {
    "customer": [{
      "customerid": "BLAUS",
      "contactname": "Hanna Moos",
      "companyname": "Blauer See Delikatessen",
      "city": "Mannheim",
      "country": "Germany",
      "phone": "0621-08460"
    }],
    "orders": [{
      "orderid": 10501,
      "orderdate": "1995-05-10",
      "shipcountry": "Germany",
      "shipper": "Federal Shipping",
      "total": 149.0000000
    }, {
      "orderid": 10853,
      "orderdate": "1996-02-27",
      "shipcountry": "Germany",
      "shipper": "United Package",
      "total": 625.0000000
    }, {
      "orderid": 10956,
      "orderdate": "1996-04-16",
      "shipcountry": "Germany",
      "shipper": "United Package",
      "total": 677.0000000
    }, {
      "orderid": 11058,
      "orderdate": "1996-05-29",
      "shipcountry": "Germany",
      "shipper": "Federal Shipping",
      "total": 858.0000000
    }
  ]
}
```

# CPPServer

The complete specification of the JSON output supported by CPPServer:

[https://cppserver.com/docs/json\\_response\\_spec.pdf](https://cppserver.com/docs/json_response_spec.pdf)

## The *modern* C++ Backend

Are we going to write C++ code to create the microservice? No, we don't need to write any code, we'll just declare our intention in a JSON configuration file, CPPServer program loads this configuration and executes its generic functions that interact with the database server and produce JSON output at very high speed.

The JSON configuration for this particular example looks like this:

```
{
  "db": "demodb",
  "uri": "/ms/customer/info",
  "sql": "execute sp_getCustomerInfo($customerid)",
  "function": "dbgetm",
  "tags": [ {"tag": "customer"}, {"tag": "orders"} ],
  "fields": [
    {"name": "customerid", "type": "string", "required": "true"}
  ]
}
```

The configuration contains the *URI* for this service, the *SQL* to execute the stored procedure shown in page 1, the *\$customerid* is a placeholder that will be replaced by the corresponding input parameter, CPPServer ensures that there is no SQL injection attack in the input parameters, the *tags* array indicates the name of each JSON array (SQL resultset), and the *fields* array indicates the input fields required by this microservice in the URL, via GET or POST, including validation information such as data type and if it's required or not. With this configuration CPPServer will do all the work for you. The *function* attribute is the name to identify the internal CPPServer generic function (the actual microservice) that will be used for this case, "dbgetm" stands for "get multiple resultsets", there is also "dbget" for single-resultset cases. Internally, CPPServer will have a pointer to the function for a fast thread-safe invocation.

This JSON definition is loaded and parsed once only and transformed into a memory data structure for very fast lookups.

# CPPServer

The actual C++ code, which you don't have to write, is actually very easy to write, thanks to the efficient high-level abstractions of C++ and CPPServer:

```
//returns a single resultset
void dbget(std::string& jsonBuffer, config::microService& ms)
{
    sql::get_json(ms.db, jsonBuffer, ms.reqParams.sql( ms.sql ));
}
```

This is the native C++ code, one line of code, the sql:: component is one of the main CPPServer abstractions, encapsulates the database server native API in a powerful and pragmatic way. Each HTTP request runs on a separate thread, by using C++ *thread\_local* variables the threads share almost nothing to minimize concurrency bottlenecks. CPPServer uses a Linux-specific network server model, based on EPOLL API, it's an event oriented, async server, with one thread it can serve thousands of concurrent requests, a small pool of worker threads do the database I/O processing in the background, this way CPPServer consumes very little resources, even under heavy load.

The function will return a JSON response according to the specification, it can be OK, INVALID or ERROR response, the microservice does not care about it, detailed error description will be printed to STDERR in case of error.

The C++ code was written to minimize allocations on each request by pre-allocating buffers for json output and database connections only once when the server starts.

It's not hard to write a C++ microservice using CPPServer facilities, but for the most common tasks this is not required, you can achieve instant productivity using the declarative way and enjoy the benefits of highly-optimized native machine code instead of using interpreted languages for your server-side applications. Queries, data modification (insert/update/delete) and blob upload/download, are all covered by the pre-built microservices.

## Testing the microservice

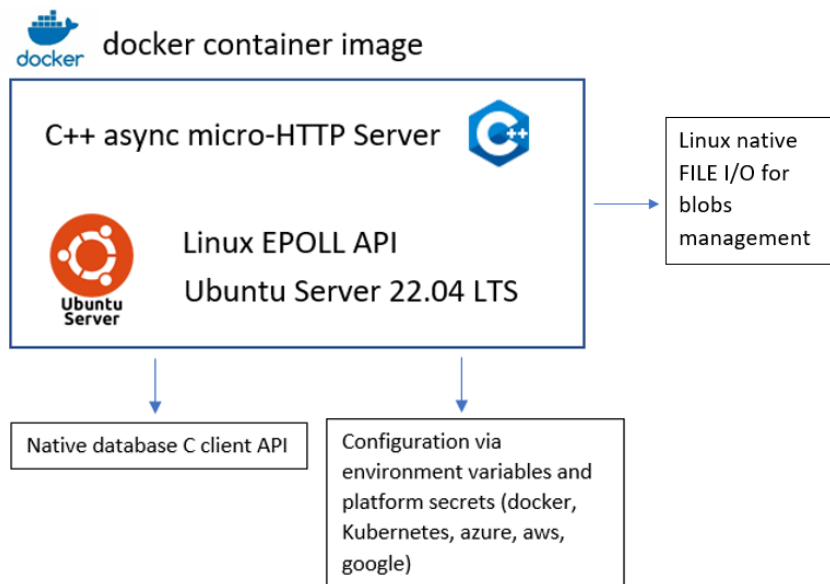
It's just a matter of sending a GET request using a browser or a command line utility such as CURL (available on Linux and Windows):

<https://cppserver.com/ms/customer/info?customerid=BERGS>

You can test other customer IDs like: ALFKI, ERNSH, BOTTM, QUICK, ANATR... also invalid input or no input at all. There will be a JSON response for every case.

The execution of the microservice is subject to security restrictions, like an existing security session started by a login (LDAP/custom DB security model) and specific roles constraints, the restrictions can be disabled for development and demos, as in the example above. If you have a database for managing logins, we can build an efficient CPPServer login adapter using your DBMS native API, a generic and configurable LDAP adapter that uses the standard LDAP native API is available too, works with OpenLDAP and ActiveDirectory.

## From laptop to cloud



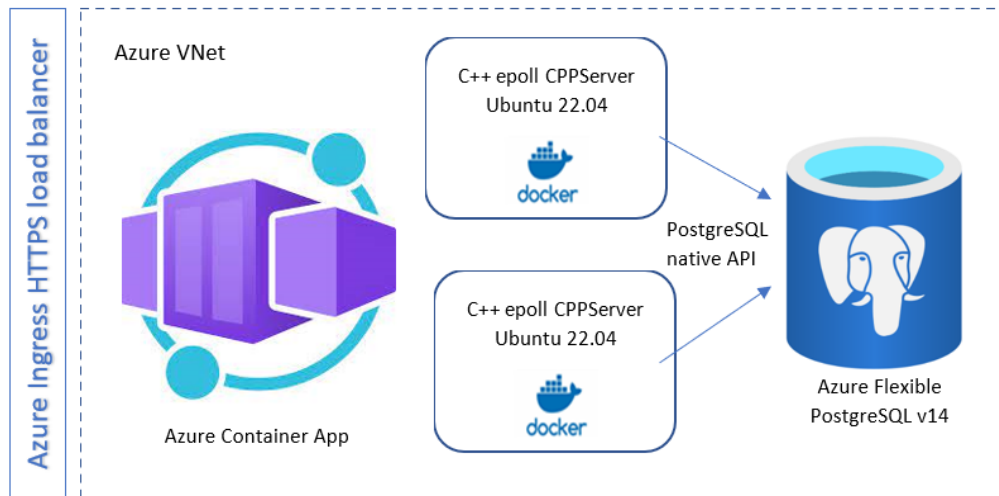
CPPServer was designed and built for Kubernetes and Cloud, running as a container using a docker image, it can be used on single-node Clusters (for testing and development) or on high-availability clusters. It's vendor agnostic, not tied to any Cloud-specific service.

# CPPServer

## Native Code

CPPServer is written in C++, compiled to platform-specific optimized machine code using GNU G++ v12.1 on Ubuntu 22.04, its executable weights about 320K.

It runs as a Pod/Container in cluster mode behind an ingress/load balancer, it's stateless, no need for session affinity. The HTTPS/TLS protocol is managed by the load balancer, the connection to the backend containers is plain HTTP over a closed network to avoid overhead, the containers are not visible outside this network, only the load balancer can reach them.



## The C++ promise

Direct mapping to hardware, efficient high-level abstractions, type-safety and resource management, we built over these fundamental benefits of using *Modern C++* and achieved the goal of having the productivity and ease of use of platforms like NodeJS or Python with the raw power of C/C++ and the Linux Kernel. CPPServer is free to use. Even for commercial purposes, commercial support is available.

QuickStart in 10 minutes:

<https://github.com/cppservergit/cppserver-docs/blob/main/quickstart.md>